

Practical Assessment (Assignment) Even Semester

Course: B.Sc. (H) Computer Science Semester VI (Section B)

Paper Name: Data Mining

Date: 29th April 2023

Name of the Candidate: Vishal Maurya

Roll No.: 5750

Examination Roll No.: 20035570117

Mobile Number: 9354643132

e-Mail Id: vishal205750@keshav.du.ac.in

Submitted To: Nidhi Ma'am

Q1. Create a file “people.txt” with the following data:

Age	agegroup	height	status	yearsmarried
21	adult	6.0	single	-1
2	child	3	married	0
18	adult	5.7	married	20
221	elderly	5	widowed	2
34	child	-7	married	3

i) Read the data from the file “people.txt”.

```
#i). Read the data from the file "people.txt"
```

```
df = pd.read_table("1stQuestionData.txt")
df
```

```
   Age  agegroup  height  status  yearsmarried
0   21    adult    6.0    single           -1
1    2    child    3.0    married            0
2   18    adult    5.7    married           20
3  221   elderly    5.0  widowed            2
4   34    child   -7.0    married            3
```

```
df.isnull().sum()
```

```
Age          0
agegroup     0
height       0
status       0
yearsmarried 0
dtype: int64
```

```
df.describe()
```

```
           Age    height  yearsmarried
count  5.000000  5.000000    5.000000
mean   59.200000  2.540000    4.800000
std    91.163041  5.45967    8.642916
min     2.000000 -7.000000   -1.000000
25%    18.000000  3.000000    0.000000
50%    21.000000  5.000000    2.000000
75%    34.000000  5.700000    3.000000
max   221.000000  6.000000   20.000000
```

ii) Create a ruleset E that contain rules to check for the following conditions:

1. The age should be in the range 0-150.
2. The age should be greater than years married.
3. The status should be married or single or widowed.
4. If age is less than 18 the age group should be child, if age is between 18 and 65 the age group should be adult, if age is more than 65 the age group should be elderly.

```
# ii). Create a ruleset E that contain rules to check for the following conditions:

"""
1. The age should be in the range 0-150.
2. The age should be greater than yearsmarried.
3. The status should be married or single or widowed.
4. If age is less than 18 the agegroup should be child, if age is between 18 and 65 the agegroup should be adult,
   if age is more than 65 the agegroup should be elderly.
"""

def ruleset(df):
    df['Rule1'] = df['Age'].apply(lambda x: True if x > 0 and x < 150 else False)
    df['Rule2'] = df.apply(lambda x: True if x.Age > x.yearsmarried else False, axis = 1)
    df['Rule3'] = df['status'].apply(lambda x: True if x == 'married' or x == 'single' or x == 'widowed' else False)
    df['Rule4'] = df.apply(lambda x: True if ((x.Age < 18 and x.agegroup == 'child') or
                                             (x.Age >= 18 and x.Age <= 65 and x.agegroup == 'adult') or
                                             (x.Age > 65 and x.agegroup == 'elderly'))
                           else False , axis = 1)
```

iii) Check whether ruleset E is violated by the data in the file people.txt.

```
# iii). Check whether ruleset E is violated by the data in the file people.txt.
```

```
ruleset(df)
df
```

	Age	agegroup	height	status	yearsmarried	Rule1	Rule2	Rule3	Rule4
0	21	adult	6.0	single	-1	True	True	True	True
1	2	child	3.0	married	0	True	True	True	True
2	18	adult	5.7	married	20	True	False	True	True
3	221	elderly	5.0	widowed	2	False	True	True	True
4	34	child	-7.0	married	3	True	True	True	False

iv) Summarize the results obtained in part (iii).

```
df_rule_followed = df.iloc[:, 5:]

df_rule_followed = df_rule_followed.astype(int)

df_rule_followed    # here 1 represents that rule is followed and 0 represents that rule is violated
```


	Rule1	Rule2	Rule3	Rule4
0	1	1	1	1
1	1	1	1	1
2	1	0	1	1
3	0	1	1	1
4	1	1	1	0

```
df_rule_followed.describe()
```

	Rule1	Rule2	Rule3	Rule4
count	5.000000	5.000000	5.0	5.000000
mean	0.800000	0.800000	1.0	0.800000
std	0.447214	0.447214	0.0	0.447214
min	0.000000	0.000000	1.0	0.000000
25%	1.000000	1.000000	1.0	1.000000
50%	1.000000	1.000000	1.0	1.000000
75%	1.000000	1.000000	1.0	1.000000
max	1.000000	1.000000	1.0	1.000000

v) Visualize the results obtained in part (iii)

```
plt.figure()
df_rule_followed.apply(lambda x:len(x) - x.sum()).plot(kind='bar')
plt.xlabel('Rules')
plt.ylabel('Number of rules violated')
```



Q2. Perform the following preprocessing tasks on the dirty_iris dataset.

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
```

```
df=pd.read_csv("dirty_iris.csv")
```

```
df.shape
```

```
(150, 5)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Sepal.Length    140 non-null   float64
1   Sepal.Width     133 non-null   float64
2   Petal.Length    131 non-null   float64
3   Petal.Width     138 non-null   float64
4   Species         150 non-null   object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

```
df.describe()
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
count	140.000000	133.000000	131.000000	138.0
mean	6.559286	3.390977	4.449962	inf
std	6.800940	3.315310	5.769299	NaN
min	0.000000	-3.000000	0.000000	0.1
25%	5.100000	2.800000	1.600000	0.3
50%	5.750000	3.000000	4.500000	1.3
75%	6.400000	3.300000	5.100000	1.8
max	73.000000	30.000000	63.000000	inf

```
df.head()
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
0	6.4	3.2	4.5	1.5	versicolor
1	6.3	3.3	6.0	2.5	virginica
2	6.2	NaN	5.4	2.3	virginica
3	5.0	3.4	1.6	0.4	setosa
4	5.7	2.6	3.5	1.0	versicolor

i) Calculate the number and percentage of observations that are complete.

```
# i) Calculate the number and percentage of observations that are complete.
df.isna().sum()
```

```
Sepal.Length    10
Sepal.Width      17
Petal.Length     19
Petal.Width      12
Species           0
dtype: int64
```

```
df1 = df.dropna(thresh = 5)
```

```
df1.head()
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
0	6.4	3.2	4.5	1.5	versicolor
1	6.3	3.3	6.0	2.5	virginica
3	5.0	3.4	1.6	0.4	setosa
4	5.7	2.6	3.5	1.0	versicolor
7	5.9	3.0	5.1	1.8	virginica

```
num1 = df1.shape
num = df.shape
```

```
num1 = list(num1)
num = list(num)
print("Number of observations that are complete : ", num1[0])
print("Percentage of observations that are complete : ", (num1[0]/num[0])*100)
```

```
Number of observations that are complete : 96
Percentage of observations that are complete : 64.0
```

ii) Replace all the special values in data with NA.

```
# ii) replace all the special values in data with NA.
```

```
df.iloc[:,4] = df.iloc[:,4].replace(np.inf, np.nan)  
df.iloc[:,4] = df.iloc[:,4].replace([np.nan, np.nan])
```

```
df.head()
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
0	6.4	3.2	4.5	1.5	versicolor
1	6.3	3.3	6.0	2.5	virginica
2	6.2	0.0	5.4	2.3	virginica
3	5.0	3.4	1.6	0.4	setosa
4	5.7	2.6	3.5	1.0	versicolor

iii) Define these rules in a separate text file and read them. (Use editfile function in R (package editrules). Use similar function in Python).

Print the resulting constraint object.

- Species should be one of the following values: setosa, versicolor or virginica.
- All measured numerical properties of an iris should be positive.
- The petal length of an iris is at least 2 times its petal width.
- The sepal length of an iris cannot exceed 30 cm.
- The sepals of an iris are longer than its petals.

```
# iii) Define these rules in a separate text file and read them.
```

```
# """
```

```
#     Print the resulting constraint object.
```

```
#     - Species should be one of the following values: setosa, versicolort, virginica.
```

```
#     - All measured numerical properties of an iris shouls be positive
```

```
#     - The petal length of an iris is at least 2 rimes its petal width
```

```
#     - The sepal length of an iris cannot exceed 30 cm
```

```
#     - The sepals of an iris are longer than its petals.
```

```
# """
```

```
rule1 = df['Species'].apply(lambda x:True if x=='setosa' or x=="versicolor" or x=='virginica' else False)
```

```
rule2 = df.iloc[:, :4].apply(lambda x: True if all(y > 0 for y in x) else False, axis=1)
```

```
rule3 = df.apply(lambda x:True if x["Petal.Length"]>=x['Petal.Width'] else False, axis=1)
```

```
rule4 = df['Sepal.Length'].apply(lambda x:True if x<=30 else False)
```

```
rule5 = df.apply(lambda x:True if x["Sepal.Length"]>x['Petal.Length'] else False, axis=1)
```

<pre>rule1 0 True 1 True 2 True 3 True 4 True ... 145 True 146 True 147 True 148 True 149 True Name: Species, Length: 150, dtype: bool</pre>	<pre>rule3 0 True 1 True 2 True 3 True 4 True ... 145 True 146 True 147 True 148 False 149 True Length: 150, dtype: bool</pre>
<pre>rule2 0 True 1 True 2 False 3 True 4 True ... 145 True 146 True 147 True 148 False 149 False Length: 150, dtype: bool</pre>	<pre>rule4 0 True 1 True 2 True 3 True 4 True ... 145 True 146 True 147 True 148 True 149 True Name: Sepal.Length, Length: 150, dtype: bool</pre>
<pre>rule5 0 True 1 True 2 True 3 True 4 True ... 145 True 146 True 147 True 148 True 149 True Length: 150, dtype: bool</pre>	

iv) Determine how often each rule is broken (violatedEdits). Also summarize and plot the result.


```
# iv) Determine how often each rule is broken (violatedEdits). Also summarize and plot the result.

df_rules = pd.DataFrame({"Rule1":rule1, "Rule2":rule2, "Rule3":rule3, "Rule4":rule4, "Rule5":rule5})
df_rules = df_rules.astype(int)
df_rules
```

	Rule1	Rule2	Rule3	Rule4	Rule5
0	1	1	1	1	1
1	1	1	1	1	1
2	1	0	1	1	1
3	1	1	1	1	1
4	1	1	1	1	1
...
145	1	1	1	1	1
146	1	1	1	1	1
147	1	1	1	1	1
148	1	0	0	1	1
149	1	0	1	1	1

150 rows x 5 columns

```
df_rules["Rule1"].value_counts()
```

```
1    150
Name: Rule1, dtype: int64
```

```
df_rules["Rule2"].value_counts()
```

```
1     92
0     58
Name: Rule2, dtype: int64
```

```
df_rules["Rule3"].value_counts()
```

```
1    129
0     21
Name: Rule3, dtype: int64
```

```
df_rules["Rule4"].value_counts()
```

```
1    148
0      2
Name: Rule4, dtype: int64
```


```
df_rules["Rule5"].value_counts()
```

```
1    138
0     12
Name: Rule5, dtype: int64
```

```
plt.figure()

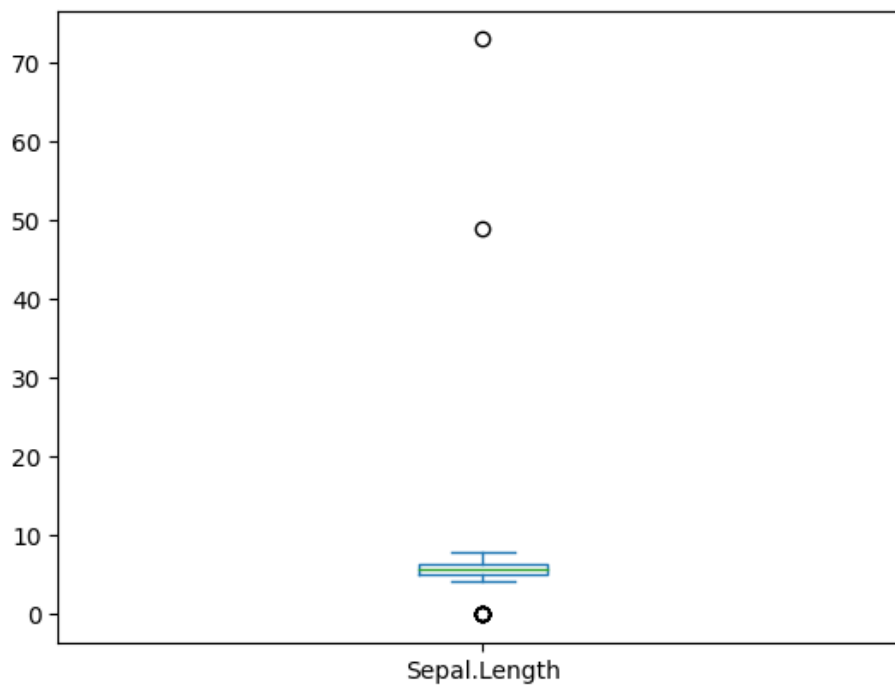
df_rules.apply(lambda x:len(x) - x.sum()).plot(kind='bar')

plt.xlabel("Rules")
plt.ylabel("Number of records that violtes the rule")
```



v) Find outliers in sepal length using boxplot and boxplot.stats

```
#v). Find outliers in sepal length using boxplot and boxplot.stats  
df["Sepal.Length"].plot(kind="box")
```



```
quantile = df["Sepal.Length"].quantile([0.0, 0.25, 0.5, 0.75, 1])
```

```
quantile
```

```
0.00    0.0
0.25    5.0
0.50    5.7
0.75    6.4
1.00   73.0
Name: Sepal.Length, dtype: float64
```

```
df["Sepal.Length"].describe()
```

```
count    150.000000
mean      6.122000
std       6.770791
min       0.000000
25%       5.000000
50%       5.700000
75%       6.400000
max      73.000000
Name: Sepal.Length, dtype: float64
```

Q3. Load the data from wine dataset. Check whether all attributes are standardized or not (mean is 0 and standard deviation is 1). If not, standardize the attributes. Do the same with Iris dataset.

```
# Load the data from wine dataset. Check whether all attributes are standardized or not (mean is 0 and standard deviation is 1).
# If not, standardize the attributes. Do the same with Iris dataset.
```

```
import pandas as pd
from sklearn.datasets import load_iris
from sklearn.datasets import load_wine
```

```
iris = load_iris()
df = pd.DataFrame(iris.data, columns=iris.feature_names)
df["Class"] = iris.target
```

For Iris Dataset:

```
df.head()
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	Class
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0

```
len(df.columns)

5

X = df.iloc[:, :4]

print("Mean : ", X.mean())
print("Standard deviation: ", X.std())

Mean : sepal length (cm)    5.843333
sepal width (cm)          3.057333
petal length (cm)         3.758000
petal width (cm)          1.199333
dtype: float64
Standard deviation: sepal length (cm)    0.828066
sepal width (cm)          0.435866
petal length (cm)         1.765298
petal width (cm)          0.762238
dtype: float64
```

```
# standardizing Iris dataset

scaled_iris_data = scaler.fit_transform(X)
scaled_iris_data = pd.DataFrame(scaled_iris_data, columns=X.columns)
scaled_iris_data.head()

   sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)
0         -0.900681         1.019004        -1.340227        -1.315444
1         -1.143017        -0.131979        -1.340227        -1.315444
2         -1.385353         0.328414        -1.397064        -1.315444
3         -1.506521         0.098217        -1.283389        -1.315444
4         -1.021849         1.249201        -1.340227        -1.315444
```

```
print("Mean : ", scaled_iris_data.mean())
print("Standard Deviation : ", scaled_iris_data.std())

Mean : sepal length (cm)    -1.690315e-15
sepal width (cm)           -1.842970e-15
petal length (cm)          -1.698641e-15
petal width (cm)           -1.409243e-15
dtype: float64
Standard Deviation : sepal length (cm)    1.00335
sepal width (cm)       1.00335
petal length (cm)       1.00335
petal width (cm)        1.00335
dtype: float64
```

For Wine Dataset

```
wine = load_wine()
df1 = pd.DataFrame(wine.data, columns=wine.feature_names)
df1["class"] = wine.target
```

```
df1.head()
```

	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium	total_phenols	flavanoids	nonflavanoid_phenols	proanthocyanins	color_intensity	hue	od280/od315_of
0	14.23	1.71	2.43		15.6	127.0	2.80	3.06	0.28	2.29	5.64	1.04
1	13.20	1.78	2.14		11.2	100.0	2.65	2.76	0.26	1.28	4.38	1.05
2	13.16	2.36	2.67		18.6	101.0	2.80	3.24	0.30	2.81	5.68	1.03
3	14.37	1.95	2.50		16.8	113.0	3.85	3.49	0.24	2.18	7.80	0.86
4	13.24	2.59	2.87		21.0	118.0	2.80	2.69	0.39	1.82	4.32	1.04

```
len(df1.columns)
```

```
14
```

```
X1 = df1.iloc[:, :13]

print("Mean : ", X1.mean())
print("Standard Deviation : ", X1.std())
```

```
Mean : alcohol          13.000618
malic_acid           2.336348
ash                 2.366517
alcalinity_of_ash    19.494944
magnesium           99.741573
total_phenols        2.295112
flavanoids          2.029270
nonflavanoid_phenols 0.361854
proanthocyanins      1.590899
color_intensity      5.058090
hue                 0.957449
od280/od315_of_diluted_wines 2.611685
proline             746.893258
dtype: float64
Standard Deviation : alcohol          0.811827
malic_acid           1.117146
ash                 0.274344
alcalinity_of_ash    3.339564
magnesium           14.282484
total_phenols        0.625851
flavanoids          0.998859
nonflavanoid_phenols 0.124453
proanthocyanins      0.572359
color_intensity      2.318286
hue                 0.228572
od280/od315_of_diluted_wines 0.709990
proline             314.907474
dtype: float64
```

```
# standardizing wine dataset
```

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()

scaled_data = scaler.fit_transform(X1)
```

```
scaled_wine_data = pd.DataFrame(scaled_data, columns=X1.columns)
scaled_wine_data.head()
```

	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium	total_phenols	flavanoids	nonflavanoid_phenols	proanthocyanins	color_intensity	hue	od28
0	1.518613	-0.562250	0.232053	-1.169593	1.913905	0.808997	1.034819	-0.659563	1.224884	0.251717	0.362177	
1	0.246290	-0.499413	-0.827996	-2.490847	0.018145	0.568648	0.733629	-0.820719	-0.544721	-0.293321	0.406051	
2	0.196879	0.021231	1.109334	-0.268738	0.088358	0.808997	1.215533	-0.498407	2.135968	0.269020	0.318304	
3	1.691550	-0.346811	0.487926	-0.809251	0.930918	2.491446	1.466525	-0.981875	1.032155	1.186068	-0.427544	
4	0.295700	0.227694	1.840403	0.451946	1.281985	0.808997	0.663351	0.226796	0.401404	-0.319276	0.362177	

```
print("Mean :", scaled_wine_data.mean())
print("Standard Deviation :", scaled_wine_data.std())
```

```
Mean : alcohol          7.841418e-15
malic_acid          2.444986e-16
ash          -4.059175e-15
alcalinity_of_ash    -7.110417e-17
magnesium          -2.494883e-17
total_phenols        -1.955365e-16
flavanoids          9.443133e-16
nonflavanoid_phenols -4.178929e-16
proanthocyanins      -1.540590e-15
color_intensity      -4.129032e-16
hue          1.398382e-15
od280/od315_of_diluted_wines 2.126888e-15
proline          -6.985673e-17
dtype: float64
Standard Deviation : alcohol          1.002821
malic_acid          1.002821
ash          1.002821
alcalinity_of_ash    1.002821
magnesium          1.002821
total_phenols        1.002821
flavanoids          1.002821
nonflavanoid_phenols 1.002821
proanthocyanins      1.002821
color_intensity      1.002821
hue          1.002821
od280/od315_of_diluted_wines 1.002821
proline          1.002821
dtype: float64
```

Q4. Run Apriori algorithm to find frequent itemsets and association rules

```
!pip install efficient_apriori
```

```
Requirement already satisfied: efficient_apriori in c:\users\admin\anaconda3\lib\site-packages (2.0.2)
```

```
from efficient_apriori import apriori
import pandas as pd
```

```
df=pd.read_csv("groceriesDataset.csv")
```

```
df.head()
```

	Item(s)	Item 1	Item 2	Item 3	Item 4	Item 5	Item 6	Item 7	Item 8	Item 9	...	Item 23	Item 24	Item 25	Item 26	Item 27	Item 28	Item 29	Item 30	Item 31	Item 32
0	4	citrus fruit	semi-finished bread	margarine	ready soups	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
1	3	tropical fruit	yogurt	coffee	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2	1	whole milk	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
3	4	pip fruit	yogurt	cream cheese	meat spreads	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
4	4	other vegetables	whole milk	condensed milk	long life bakery product	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

5 rows x 33 columns

```
df.shape
```

```
(9835, 33)
```

```
df["Item(s)"].value_counts()
```

```
1      2159
2      1643
3      1299
4      1005
5       855
6       645
7       545
8       438
9       350
10      246
11      182
12      117
13       78
14       77
15       55
16       46
17       29
19       14
18       14
21       11
20        9
23        6
22        4
29        3
26        1
32        1
27        1
```

Item(s)	0
Item 1	0
Item 2	2159
Item 3	3802
Item 4	5101
Item 5	6106
Item 6	6961
Item 7	7606
Item 8	8151
Item 9	8589
Item 10	8939
Item 11	9185
Item 12	9367
Item 13	9484
Item 14	9562
Item 15	9639
Item 16	9694
Item 17	9740
Item 18	9769
Item 19	9783
Item 20	9797
Item 21	9806
Item 22	9817
Item 23	9821
Item 24	9827
Item 25	9828
Item 26	9828
Item 27	9829
Item 28	9830
Item 29	9831
Item 30	9834
Item 31	9834
Item 32	9834

```
[ ] df.dropna(thresh=5, inplace=True)
```

 `df.head()`


```
df.isnull().sum()
```

```
Item(s)      0
Item 1       0
Item 2       0
Item 3       0
Item 4       0
Item 5     1005
Item 6     1860
Item 7     2505
Item 8     3050
Item 9     3488
Item 10     3838
Item 11     4084
Item 12     4266
Item 13     4383
Item 14     4461
Item 15     4538
Item 16     4593
Item 17     4639
Item 18     4668
Item 19     4682
Item 20     4696
Item 21     4705
Item 22     4716
Item 23     4720
Item 24     4726
Item 25     4727
Item 26     4727
Item 27     4728
Item 28     4729
Item 29     4730
Item 30     4733
Item 31     4733
Item 32     4733
```

```
[ ] 4734 - 1005
```

```
3729
```

Drop all the columns that have less than 3729 non-nan values

```
[ ] df.dropna(thresh=3729, inplace=True, axis=1)
```

```
df.head()
```

	Item(s)	Item 1	Item 2	Item 3	Item 4	Item 5
0	4	citrus fruit	semi-finished bread	margarine	ready soups	NaN
3	4	pip fruit	yogurt	cream cheese	meat spreads	NaN
4	4	other vegetables	whole milk	condensed milk	long life bakery product	NaN
5	5	whole milk	butter	yogurt	rice	abrasive cleaner
7	5	other vegetables	UHT-milk	rolls/buns	bottled beer	liquor (appetizer)

```
[ ] df.isnull().sum()
```

```
Item(s)      0
Item 1       0
Item 2       0
Item 3       0
Item 4       0
Item 5     1005
dtype: int64
```

```
df = df.fillna(method="ffill", limit=3)
```

```
df.isnull().sum()
```

```
Item(s)      0
Item 1       0
Item 2       0
Item 3       0
Item 4       0
Item 5      14
dtype: int64
```

```
df = df.fillna(method="bfill", limit=3)
```

```
df.isnull().sum()
```

```
Item(s)      0
Item 1       0
Item 2       0
Item 3       0
Item 4       0
Item 5       0
dtype: int64
```

```
df.drop(["Item(s)", axis=1, inplace=True)
```

```
df1=df.values.tolist()
```

4.1 Use minimum support as 2% and minimum confidence as 5%

Taking minimum support threshold 0.02 and minimum confidence threshold 0.05

```
[ ] itemsets, rules = apriori(df1, min_support=0.02, min_confidence=0.05)
```

```
[ ] itemsets
```

```
{1: {'citrus fruit',): 691,
      ('margarine',): 275,
      ('pip fruit',): 609,
      ('yogurt',): 913,
      ('cream cheese',): 232,
      ('other vegetables',): 1497,
      ('whole milk',): 1792,
      ('long life bakery product',): 117,
      ('butter',): 406,
      ('UHT-milk',): 236,
      ('rolls/buns',): 886,
      ('bottled beer',): 256,
      ('tropical fruit',): 866,
      ('white bread',): 197,
      ('bottled water',): 459,
      ('chocolate',): 131,
      ('curd',): 385,
      ('sugar',): 132,
      ('fruit/vegetable juice',): 230,
      ('newspapers',): 192,
      ('processed cheese',): 98,
      ('root vegetables',): 915,
      ('sausage',): 723,
      ('soda',): 644,
      ('brown bread',): 324,
      ('canned beer',): 198,
      ('beverages',): 154,
```

```
[{other vegetables} -> {beef},
{beef} -> {other vegetables},
{root vegetables} -> {beef},
{beef} -> {root vegetables},
{whole milk} -> {beef},
{beef} -> {whole milk},
{soda} -> {bottled water},
{bottled water} -> {soda},
{whole milk} -> {bottled water},
{bottled water} -> {whole milk},
{whole milk} -> {brown bread},
{brown bread} -> {whole milk},
{other vegetables} -> {butter},
{butter} -> {other vegetables},
{whole milk} -> {butter},
{butter} -> {whole milk},
{other vegetables} -> {chicken},
{chicken} -> {other vegetables},
{whole milk} -> {chicken},
{chicken} -> {whole milk},
{other vegetables} -> {citrus fruit},
{citrus fruit} -> {other vegetables},
{pip fruit} -> {citrus fruit},
{citrus fruit} -> {pip fruit},
{root vegetables} -> {citrus fruit},
{citrus fruit} -> {root vegetables},
{sausage} -> {citrus fruit},
{citrus fruit} -> {sausage},
{tropical fruit} -> {citrus fruit},
```

Taking minimum support threshold 0.1 and minimum confidence threshold 0.2

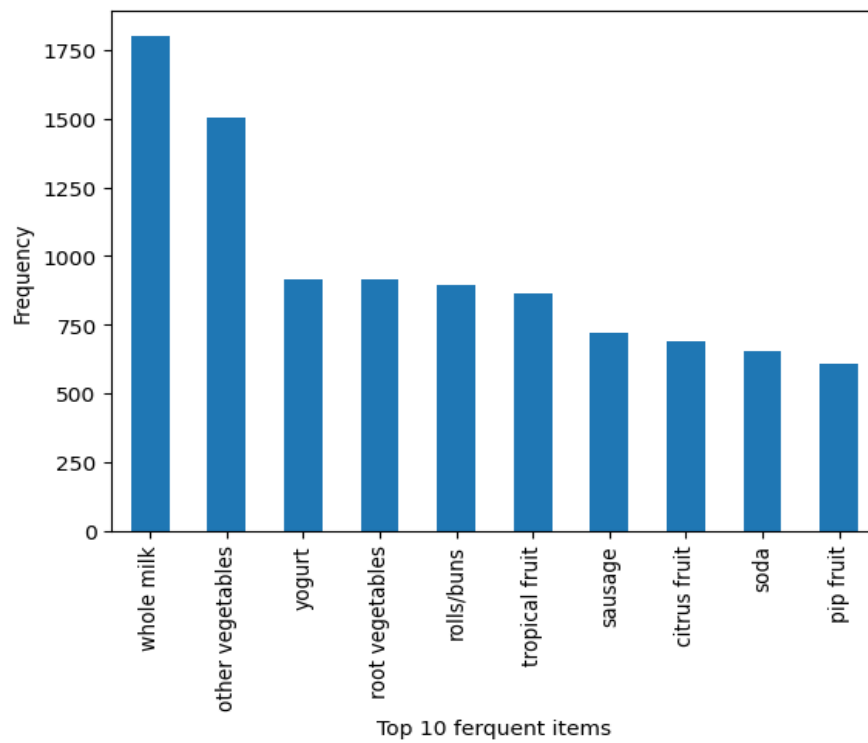
Plotting the most frequent items

```
[29] items = df.stack()

import matplotlib.pyplot as plt

plt.figure()
items.value_counts().head(10).plot(kind = "bar")

plt.xlabel("Top 10 ferquent itemset")
plt.ylabel("Frequency")
```



Another Dataset

Reading Dataset

```
[31] new_df = pd.read_csv('1000i.csv')
```

```
[32] new_df.columns = ["Receipt", "food", "quantity"]
```

```
[33] new_df.head()
```

	Receipt	food	quantity
0	1	4	15
1	1	2	49
2	1	5	44
3	2	1	1
4	2	2	19

```
new_df.isna().sum()
```

```
Receipt    0  
food       0  
quantity   0  
dtype: int64
```

```
new_df.shape
```

```
(3537, 3)
```

```
new_df.describe()
```

	Receipt	food	quantity
count	3537.000000	3537.000000	3537.000000
mean	498.145321	2.986146	24.223353
std	289.283458	1.403204	14.722389
min	1.000000	1.000000	0.000000
25%	244.000000	2.000000	12.000000
50%	496.000000	3.000000	24.000000
75%	754.000000	4.000000	37.000000
max	1000.000000	5.000000	49.000000

```
new_df["food"] = new_df["food"].map({1:"milk",2:"sugar",3:"chocolate",4:"apples",5:"curd"})
```

```
new_df.head(2)
```

	Receipt	food	quantity
0	1	apples	15
1	1	sugar	49

```
li = []
```

```
for i in new_df["Receipt"].unique():  
    li.append([])
```

```
for i in range(new_df.shape[0]):  
    li[new_df["Receipt"][i] - 1].append(new_df["food"][i])
```

```
li
```

```
[['apples', 'sugar', 'curd'],  
 ['milk', 'sugar'],  
 ['milk', 'milk'],  
 ['milk', 'milk', 'curd', 'curd', 'milk', 'milk'],  
 ['apples', 'apples', 'sugar', 'curd', 'curd'],  
 ['sugar', 'apples', 'chocolate'],  
 ['apples', 'chocolate', 'sugar', 'milk', 'milk', 'milk'],
```

Taking minimum support threshold 0.02 and minimum confidence threshold 0.05

```
[ ] itemsets, rules = apriori(li, min_support=0.02, min_confidence=0.05)
```

```
[ ] itemsets
```

```
{1: {'apples',): 540,  
 ('sugar',): 498,  
 ('curd',): 493,  
 ('milk',): 542,  
 ('chocolate',): 506},  
 2: {'apples', 'chocolate'): 250,  
 ('apples', 'curd'): 264,  
 ('apples', 'milk'): 271,  
 ('apples', 'sugar'): 258,  
 ('chocolate', 'curd'): 231,  
 ('chocolate', 'milk'): 251,  
 ('chocolate', 'sugar'): 239,  
 ('curd', 'milk'): 240,  
 ('curd', 'sugar'): 230,  
 ('milk', 'sugar'): 254},  
 3: {'apples', 'chocolate', 'curd'): 112,  
 ('apples', 'chocolate', 'milk'): 120,  
 ('apples', 'chocolate', 'sugar'): 116,  
 ('apples', 'curd', 'milk'): 120,  
 ('apples', 'curd', 'sugar'): 122,  
 ('apples', 'milk', 'sugar'): 124,  
 ('chocolate', 'curd', 'milk'): 115,  
 ('chocolate', 'curd', 'sugar'): 115,  
 ('chocolate', 'milk', 'sugar'): 115}
```

```
rules
```

```
[{chocolate} -> {apples},  
{apples} -> {chocolate},  
{curd} -> {apples},  
{apples} -> {curd},  
{milk} -> {apples},  
{apples} -> {milk},  
{sugar} -> {apples},  
{apples} -> {sugar},  
{curd} -> {chocolate},  
{chocolate} -> {curd},  
{milk} -> {chocolate},  
{chocolate} -> {milk},  
{sugar} -> {chocolate},  
{chocolate} -> {sugar},
```

Taking minimum support threshold 0.1 and minimum confidence threshold 0.2

```
[ ] itemsets, rules = apriori(df1, min_support=0.1, min_confidence=0.2)
```

```
[ ] itemsets
```

```
{1: {'citrus fruit',): 691,  
 ('pip fruit',): 609,  
 ('yogurt',): 913,  
 ('other vegetables',): 1497,  
 ('whole milk',): 1792,  
 ('rolls/buns',): 886,  
 ('tropical fruit',): 866,  
 ('root vegetables',): 915,  
 ('sausage',): 723,  
 ('soda',): 644},  
 2: {'other vegetables', 'whole milk'): 577}}
```

```
[ ] rules
```

```
[{whole milk} -> {other vegetables}, {other vegetables} -> {whole milk}]
```

Plotting the most frequent item

```
[ ] unique_food = new_df["food"].unique()

[ ] unique_food

array(['apples', 'sugar', 'curd', 'milk', 'chocolate'], dtype=object)

[ ] count = {}

for i in range(new_df.shape[0]):
    if new_df["food"][i] not in count:
        count[new_df["food"][i]] = new_df["quantity"][i]
    else:
        count[new_df["food"][i]] = new_df["quantity"][i] + count[new_df["food"][i]]

[ ] count

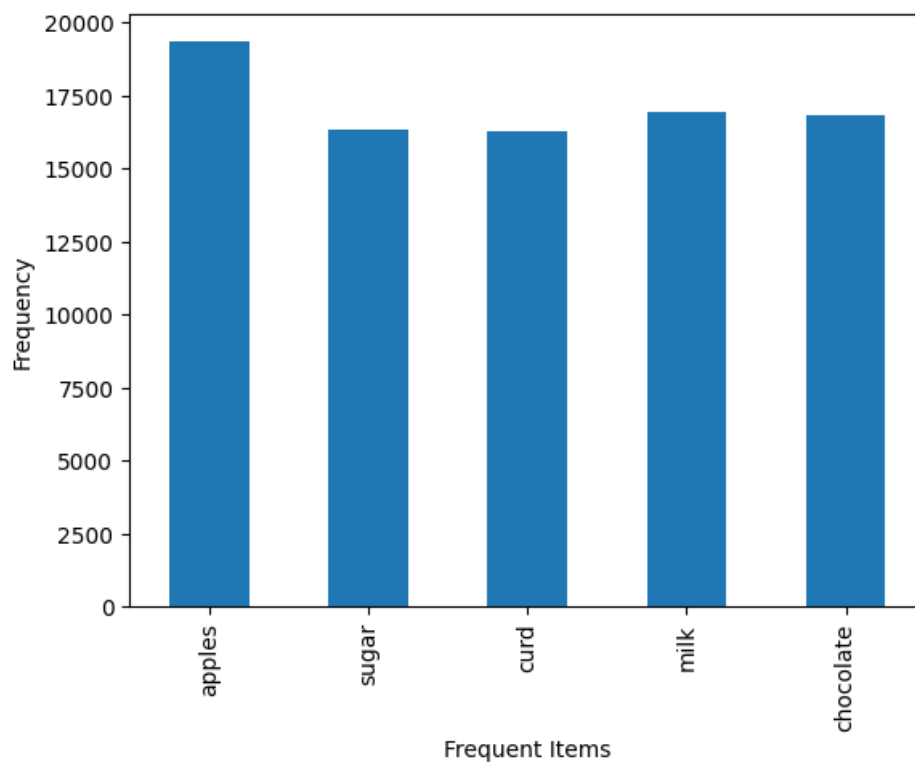
{'apples': 19337,
 'sugar': 16351,
 'curd': 16273,
 'milk': 16910,
 'chocolate': 16807}

[ ] count = pd.Series(count)
```

```
plt.figure()

count.plot(kind="bar")

plt.xlabel("Frequent Items")
plt.ylabel("Frequency")
```



Q5. Use Naive bayes, K-nearest, and Decision tree classification algorithms and build classifiers. Divide the data set into training and test set. Compare the accuracy of the different classifiers under the following situations:

5.1 a) Training set = 75% Test set = 25%

b) Training set = 66.6% (2/3rd of total), Test set = 33.3%

5.2 Training set is chosen by

i) hold out method

ii) Random subsampling

iii) Cross-Validation. Compare the accuracy of the classifiers obtained.

5.3 Data is scaled to standard format.

For Iris Dataset

Importing dataset

```
[ ] import pandas as pd
import numpy as np
from sklearn.datasets import load_iris
```

```
[ ] iris = load_iris()
```

```
[ ] X = iris.data
Y = iris.target
```

```
[ ] iris_df = pd.DataFrame(X, columns=iris.feature_names)
```

```
[ ] iris_df["class"] = Y
```

```
[ ] iris_df.head(2)
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	class
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0

```
iris_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   sepal length (cm)      150 non-null   float64
1   sepal width (cm)       150 non-null   float64
2   petal length (cm)      150 non-null   float64
3   petal width (cm)       150 non-null   float64
4   class                  150 non-null   int64
dtypes: float64(4), int64(1)
memory usage: 6.0 KB
```

```
iris_df.describe()
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	class
count	150.000000	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333	1.000000
std	0.828066	0.435866	1.765298	0.762238	0.819232
min	4.300000	2.000000	1.000000	0.100000	0.000000
25%	5.100000	2.800000	1.600000	0.300000	0.000000
50%	5.800000	3.000000	4.350000	1.300000	1.000000
75%	6.400000	3.300000	5.100000	1.800000	2.000000


```
iris_df.isnull().sum()
```

```
sepal length (cm)    0
sepal width (cm)     0
petal length (cm)    0
petal width (cm)     0
class                0
dtype: int64
```

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
```

```
scaled_iris = scaler.fit_transform(iris_df.iloc[:, :-1])
```

```
X = pd.DataFrame(scaled_iris, columns = iris.feature_names)
```

```
X.head(2)
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
0	-0.900681	1.019004	-1.340227	-1.315444
1	-1.143017	-0.131979	-1.340227	-1.315444

For Breast Cancer Dataset:

Breast cancer dataset

```
Loading...
df = pd.read_csv("/content/breast-cancer.csv")
```

```
[ ] df.head()
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	...	radius_worst	t
0	842302	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	...	25.38	
1	842517	M	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	...	24.99	
2	84300903	M	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	...	23.57	
3	84348301	M	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	...	14.91	
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	...	22.54	

5 rows x 32 columns

4

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 569 entries, 0 to 568
```

```
Data columns (total 32 columns):
```

#	Column	Non-Null Count	Dtype
0	id	569 non-null	int64
1	diagnosis	569 non-null	object
2	radius_mean	569 non-null	float64
3	texture_mean	569 non-null	float64
4	perimeter_mean	569 non-null	float64
5	area_mean	569 non-null	float64
6	smoothness_mean	569 non-null	float64
7	compactness_mean	569 non-null	float64
8	concavity_mean	569 non-null	float64
9	concave points_mean	569 non-null	float64
10	symmetry_mean	569 non-null	float64
11	fractal_dimension_mean	569 non-null	float64
12	radius_se	569 non-null	float64
13	texture_se	569 non-null	float64
14	perimeter_se	569 non-null	float64
15	area_se	569 non-null	float64
16	smoothness_se	569 non-null	float64
17	compactness_se	569 non-null	float64
18	concavity_se	569 non-null	float64
19	concave points_se	569 non-null	float64
20	symmetry_se	569 non-null	float64

```
df = df.drop(["id"], axis = 1)
```

```
df.head()
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	symmetry_mean	...	radius_wor
0	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419	...	25.
1	M	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812	...	24.
2	M	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	0.2069	...	23.
3	M	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	0.2597	...	14.
4	M	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	0.1809	...	22.

```
5 rows x 31 columns
```

```
df.describe()
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean
count	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000
mean	14.127292	19.289649	91.969033	654.889104	0.096360	0.104341	0.088799
std	3.524049	4.301036	24.298981	351.914129	0.014064	0.052813	0.079720
min	6.981000	9.710000	43.790000	143.500000	0.052630	0.019380	0.000000
25%	11.700000	16.170000	75.170000	420.300000	0.086370	0.064920	0.029560
50%	13.370000	18.840000	86.240000	551.100000	0.095870	0.092630	0.061540
75%	15.780000	21.800000	104.100000	782.700000	0.105300	0.130400	0.130700
max	28.110000	39.280000	188.500000	2501.000000	0.163400	0.345400	0.426800

```
8 rows x 30 columns
```

```
df.isnull().sum()
```

```
diagnosis      0
radius_mean    0
texture_mean    0
perimeter_mean 0
area_mean      0
smoothness_mean 0
compactness_mean 0
concavity_mean 0
concave points_mean 0
symmetry_mean  0
fractal_dimension_mean 0
radius_se      0
texture_se     0
perimeter_se   0
area_se        0
smoothness_se  0
compactness_se 0
concavity_se   0
concave points_se 0
symmetry_se    0
fractal_dimension_se 0
radius_worst   0
texture_worst  0
perimeter_worst 0
area_worst     0
smoothness_worst 0
compactness_worst 0
concavity_worst 0
concave points_worst 0
symmetry_worst 0
fractal_dimension_worst 0
dtype: int64
```

```
from sklearn.preprocessing import LabelEncoder
```

```
encoder = LabelEncoder()
```

```
df["diagnosis"] = encoder.fit_transform(df["diagnosis"])
```

```
df.head(2)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	symmetry_mean	...	radius_worst
0	1	17.99	10.38	122.8	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419	...	25.1
1	1	20.57	17.77	132.9	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812	...	24.9

```
2 rows x 13 columns
```

```
scaled_data = scaler.fit_transform(df.iloc[:, 1:])

df1 = pd.DataFrame(scaled_data, columns=list(df.columns)[1:])

df1.head(1)
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	symmetry_mean	fractal_dimension_mean	..
0	1.097064	-2.073335	1.269934	0.984375	1.568466	3.283515	2.652874	2.532475	2.217515	2.255747	

1 rows x 30 columns

```
df1["diagnosis"] = df["diagnosis"]

df1.head(1)
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	symmetry_mean	fractal_dimension_mean	..
0	1.097064	-2.073335	1.269934	0.984375	1.568466	3.283515	2.652874	2.532475	2.217515	2.255747	

1 rows x 31 columns

```
X_c = df1.drop("diagnosis", axis=1)
Y_c = df1["diagnosis"]

X_c.head(1)
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean
0	1.097064	-2.073335	1.269934	0.984375	1.568466	3.283515	2.652874

1 rows x 30 columns

```
Y_c.head()
```

```
0    1
1    1
2    1
3    1
4    1
Name: diagnosis, dtype: int64
```

Holdout Method

HOLDOUT METHOD

```
[ ] from sklearn.model_selection import train_test_split
```

Training set = 75%, test set = 25%

```
[ ] X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.25, random_state = 1)
    X_train_c, X_test_c, Y_train_c, Y_test_c = train_test_split(X_c, Y_c, test_size=0.25, random_state = 1)
```

```
[ ] print("For Iris", X_train.shape)
    print("For Breast Cancer", X_train_c.shape)
```

```
For Iris (112, 4)
For Breast Cancer (426, 30)
```

```
[ ] print("For Iris", X_test.shape)
    print("For Breast Cancer", X_test_c.shape)
```

```
For Iris (38, 4)
For Breast Cancer (143, 30)
```

```
print("For Iris", Y_train.shape)
print("For Breast Cancer", Y_train_c.shape)
```

```
For Iris (112,)
For Breast Cancer (426,)
```

```
print("For Iris", Y_test.shape)
print("For Breast Cancer", Y_test_c.shape)
```

```
For Iris (38,)
For Breast Cancer (143,)
```

Decision Tree (On Iris dataset).

```
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
```

```
dec_tree = DecisionTreeClassifier(criterion = 'entropy')
dec_tree.fit(X_train, Y_train)
prediction = dec_tree.predict(X_test)
```

```
prediction
```

```
array([0, 1, 1, 0, 2, 1, 2, 0, 0, 2, 1, 0, 2, 1, 1, 0, 1, 1, 0, 0, 1, 1,
       2, 0, 2, 1, 0, 0, 1, 2, 1, 2, 1, 2, 2, 0, 1, 0])
```

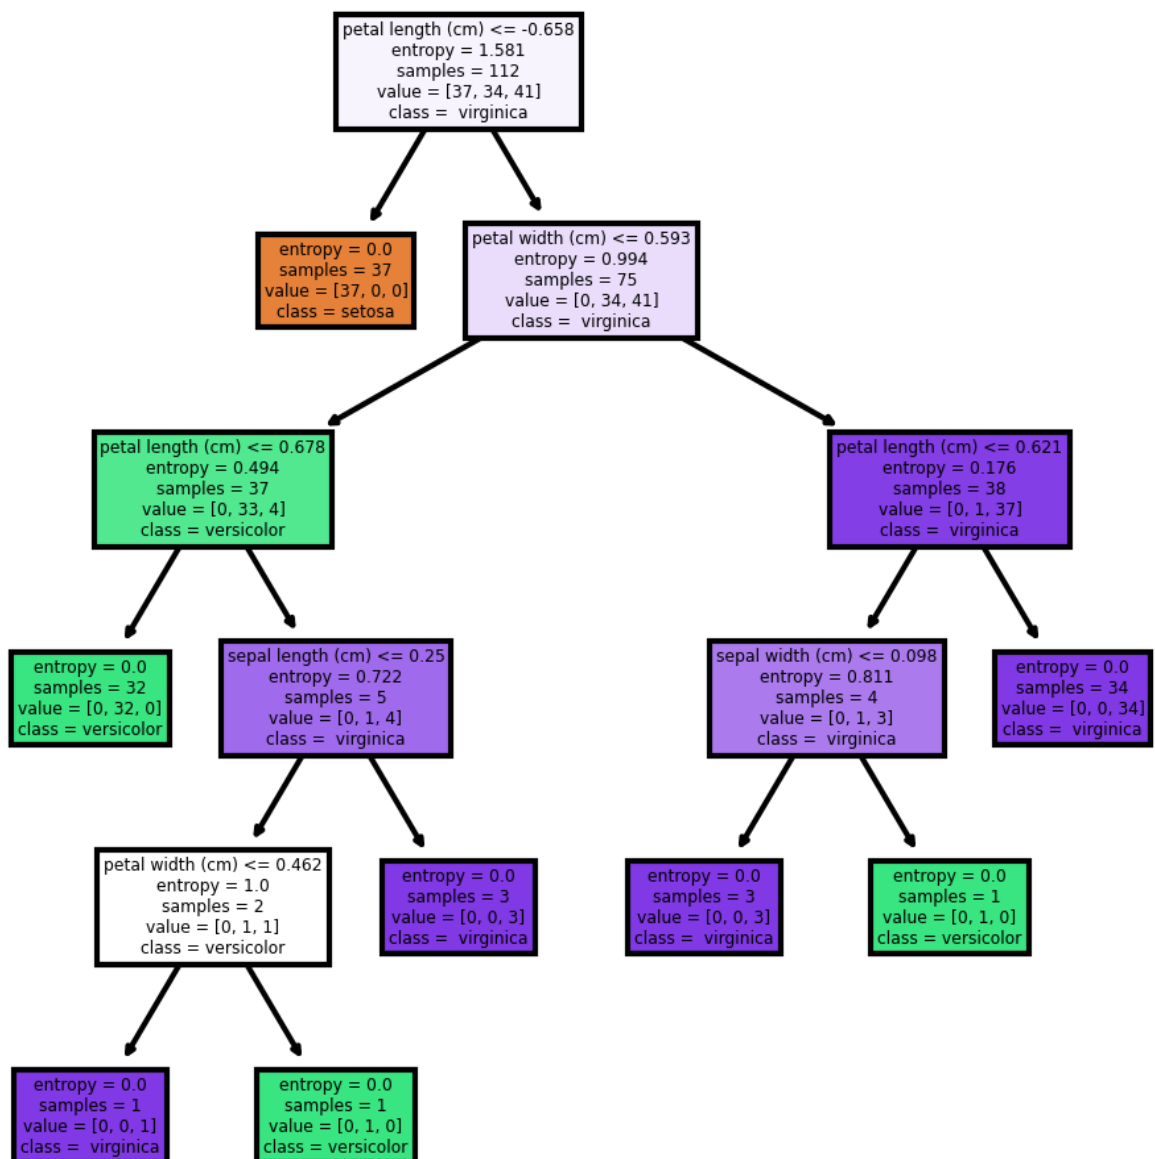
```
tree.plot_tree(dec_tree)
```

```

from matplotlib import pyplot as plt
from sklearn import metrics


fn = iris.feature_names
cn = ['setosa', 'versicolor', ' virginica']
fig, axes=plt.subplots(nrows=1,ncols=1,figsize=(4,4),dpi=300)
tree.plot_tree(dec_tree,
                feature_names=fn,
                class_names=cn,
                filled=True);

```



```
from sklearn import metrics
from sklearn.metrics import ConfusionMatrixDisplay

ConfusionMatrixDisplay.from_estimator(dec_tree, X_test, Y_test)
```



```
acc_hold_iris1 = metrics.accuracy_score(Y_test, prediction)*100
print("Accuracy on the test data ", acc_hold_iris1)
print("\nClassification report on the test data ", metrics.classification_report(Y_test, prediction))
```

Accuracy on the test data 97.36842105263158

Classification report on the test data					precision	recall	f1-score	support
0	1.00	1.00	1.00	13				
1	1.00	0.94	0.97	16				
2	0.90	1.00	0.95	9				
accuracy				38				
macro avg	0.97	0.98	0.97	38				
weighted avg	0.98	0.97	0.97	38				

Decision Tree (On Breast Cancer Dataset)

```
dec_tree1 = DecisionTreeClassifier(criterion = 'entropy')
dec_tree1.fit(X_train_c, Y_train_c)
prediction1 = dec_tree1.predict(X_test_c)

prediction1


array([0, 1, 0, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0,
       1, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0,
       0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0,
       0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0,
       1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1,
       1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 0,
       0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0])

tree.plot_tree(dec_tree1)
```


```

fn = list(df1.columns[:-1])
cn = ['malignant', 'benign']
fig, axes=plt.subplots(nrows=1,ncols=1,figsize=(6,6),dpi=300)
tree.plot_tree(dec_tree1,
                feature_names=fn,
                class_names=cn,
                filled=True);

```




```
ConfusionMatrixDisplay.from_estimator(dec_tree1, X_test_c, Y_test_c)
```



```
acc_hold_breast1 = metrics.accuracy_score(Y_test_c, prediction1)*100
print("Accuracy on the test data ", acc_hold_breast1)
print("\nClassification report on the test data ", metrics.classification_report(Y_test_c, prediction1))
```

Accuracy on the test data 94.4055944055944

Classification report on the test data					precision	recall	f1-score	support
0	0.93	0.98	0.96	88				
1	0.96	0.89	0.92	55				
accuracy			0.94	143				
macro avg	0.95	0.93	0.94	143				
weighted avg	0.94	0.94	0.94	143				

KNN (On Iris Dataset)


```
from sklearn.neighbors import KNeighborsClassifier
```

```
knn = KNeighborsClassifier()
knn.fit(X_train, Y_train)
prediction2 = knn.predict(X_test)
```

```
prediction2
```

```
array([0, 1, 1, 0, 2, 1, 2, 0, 0, 2, 1, 0, 2, 1, 1, 0, 1, 1, 0, 0, 1, 1,
       2, 0, 2, 1, 0, 0, 1, 2, 1, 2, 1, 2, 2, 0, 1, 0])
```

```
ConfusionMatrixDisplay.from_estimator(knn, X_test, Y_test)
```



```
acc_hold_iris_knn = metrics.accuracy_score(Y_test, prediction2)*100

print("Accuracy on the test data ", acc_hold_iris_knn)
print("\nClassification report on the test data ", metrics.classification_report(Y_test, prediction2))
```

Accuracy on the test data 97.36842105263158

Classification report on the test data					precision	recall	f1-score	support
0	1.00	1.00	1.00	13				
1	1.00	0.94	0.97	16				
2	0.90	1.00	0.95	9				
accuracy				0.97	38			
macro avg				0.97	0.98	0.97	38	
weighted avg				0.98	0.97	0.97	38	


KNN (On Breast Cancer Dataset)

```
knn1 = KNeighborsClassifier()
knn1.fit(X_train_c, Y_train_c)
prediction_c = knn1.predict(X_test_c)
```

prediction_c

```
array([0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0,
       1, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0,
       0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0,
       0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 1,
       1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1,
       1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 0,
       0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0])
```

```
ConfusionMatrixDisplay.from_estimator(knn1, X_test_c, Y_test_c)
```



```
acc_hold_breast_knn = metrics.accuracy_score(Y_test_c, prediction_c)*100

print("Accuracy on the test data ", acc_hold_breast_knn)
print("\nClassification report on the test data ", metrics.classification_report(Y_test_c, prediction_c))
```

Accuracy on the test data 95.1048951048951


		precision	recall	f1-score	support
0	0.94	0.99	0.96	88	
1	0.98	0.89	0.93	55	
accuracy			0.95	143	
macro avg	0.96	0.94	0.95	143	
weighted avg	0.95	0.95	0.95	143	

Naïve Bayes (On Iris Dataset)

```
from sklearn.naive_bayes import GaussianNB
```

```
nb = GaussianNB()
nb.fit(X_train,Y_train)
prediction_nb=nb.predict(X_test)
```

```
ConfusionMatrixDisplay.from_estimator(nb, X_test, Y_test)
```



```
print("Accuracy on the test data ")
acc_hold_iris_nb = metrics.accuracy_score(Y_test,prediction_nb)*100
print(f"Naive-bayes accuracy: {acc_hold_iris_nb}%")

print("\nClassification report on the test data ", metrics.classification_report(Y_test, prediction_nb))
```

```
Accuracy on the test data
Naive-bayes accuracy: 97.36842105263158%
```


```
Classification report on the test data
```

		precision	recall	f1-score	support
0	1.00	1.00	1.00	13	
1	1.00	0.94	0.97	16	
2	0.90	1.00	0.95	9	
accuracy			0.97	38	
macro avg	0.97	0.98	0.97	38	
weighted avg	0.98	0.97	0.97	38	

Naïve Bayes (On Breast Cancer Dataset)

```
nb_c = GaussianNB()
nb_c.fit(X_train_c,Y_train_c)
prediction_nbc=nb_c.predict(X_test_c)
```

```
ConfusionMatrixDisplay.from_estimator(nb_c, X_test_c, Y_test_c)
```



```
print("Accuracy on the test data ")
acc_hold_breast_nb=metrics.accuracy_score(Y_test_c,prediction_nbc)*100
print(f"Naive-bayes accuracy: {acc_hold_breast_nb}%")

print("\nClassification report on the test data ", metrics.classification_report(Y_test_c, prediction_nbc))
```

```
Accuracy on the test data
Naive-bayes accuracy: 94.4055944055944%
```

```
Classification report on the test data
```

		precision	recall	f1-score	support
0	0.94	0.97	0.96	88	
1	0.94	0.91	0.93	55	
accuracy			0.94	143	
macro avg	0.94	0.94	0.94	143	
weighted avg	0.94	0.94	0.94	143	

Random Subsampling

Decision Tree (On Iris Dataset)

```

acc = []

dec_tree = DecisionTreeClassifier(criterion = 'entropy')

for i in range(10):
    X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.25, random_state = i+1)

    dec_tree.fit(X_train, Y_train)
    prediction = dec_tree.predict(X_test)
    acc.append(metrics.accuracy_score(Y_test, prediction))

acc = sum(acc) / len(acc)

acc_rs_iris_dt = acc*100
print("Accuracy on the test data ", acc_rs_iris_dt , "%")

Accuracy on the test data  94.21052631578945 %

```

Decision Tree (Breast Cancer Dataset)

```

acc = []

dec_tree_c = DecisionTreeClassifier(criterion = 'entropy')

for i in range(10):
    X_train_c,X_test_c,Y_train_c,Y_test_c = train_test_split(X_c,Y_c,test_size=0.25, random_state = i+1)

    dec_tree_c.fit(X_train_c, Y_train_c)
    prediction_c = dec_tree_c.predict(X_test_c)
    acc.append(metrics.accuracy_score(Y_test_c, prediction_c))

acc = sum(acc) / len(acc)

acc_rs_breast_dt = acc*100
print("Accuracy on the test data ", acc_rs_breast_dt, "%")

Accuracy on the test data  93.35664335664336 %

```

KNN (Iris Dataset)

```

acc = []

knn = KNeighborsClassifier()

for i in range(10):
    X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.25, random_state = i+1)
    knn.fit(X_train, Y_train)
    prediction2 = knn.predict(X_test)
    acc.append(metrics.accuracy_score(Y_test, prediction2))

acc = sum(acc) / len(acc)

acc_rs_iris_knn = acc*100
print("Accuracy on the test data ", acc_rs_iris_knn, "%")

Accuracy on the test data  95.26315789473682 %

```

KNN (Breast Cancer Dataset)

```

acc = []

knn_c = KNeighborsClassifier()

for i in range(10):
    X_train_c,X_test_c,Y_train_c,Y_test_c = train_test_split(X_c,Y_c,test_size=0.25, random_state = i+1)
    knn_c.fit(X_train_c, Y_train_c)
    prediction2_c = knn_c.predict(X_test_c)
    acc.append(metrics.accuracy_score(Y_test_c, prediction2_c))

acc = sum(acc) / len(acc)

acc_rs_breast_knn = acc*100
print("Accuracy on the test data ", acc_rs_breast_knn, "%")

Accuracy on the test data  96.50349650349652 %

```

Naïve Bayes (On Iris Dataset)

```

acc = []

nb = GaussianNB()

for i in range(10):
    X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.25, random_state = i+1)
    nb.fit(X_train, Y_train)
    prediction_nb=nb.predict(X_test)
    acc.append(metrics.accuracy_score(Y_test, prediction_nb))

acc = sum(acc) / len(acc)

acc_rs_iris_nb = acc*100

print("Accuracy on the test data ", acc_rs_iris_nb, "%")

Accuracy on the test data  94.73684210526315 %

```

Naïve Bayes (On Breast Cancer Dataset)

```

acc = []

nb_c = GaussianNB()

for i in range(10):
    X_train_c,X_test_c,Y_train_c,Y_test_c = train_test_split(X_c,Y_c,test_size=0.25, random_state = i+1)
    nb_c.fit(X_train_c, Y_train_c)
    prediction_nbc=nb_c.predict(X_test_c)
    acc.append(metrics.accuracy_score(Y_test_c, prediction_nbc))

acc = sum(acc) / len(acc)

acc_rs_breast_nb = acc*100

print("Accuracy on the test data ", acc_rs_breast_nb, "%")

Accuracy on the test data  93.77622377622379 %

```

Training set = 66.6% (2/3rd of total), Test set = 33.3%

Holdout Method

```
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.333, random_state = 1)
X_train_c, X_test_c, Y_train_c, Y_test_c = train_test_split(X_c, Y_c, test_size=0.333, random_state = 1)
```

```
X_train.shape
```

```
(100, 4)
```

```
X_test.shape
```

```
(50, 4)
```

```
X_train_c.shape
```

```
(381, 30)
```


```
X_test_c.shape
```

```
(188, 30)
```

Decision Tree (On Iris Dataset)

```
dec_tree = DecisionTreeClassifier(criterion = 'entropy')
dec_tree.fit(X_train, Y_train)
prediction = dec_tree.predict(X_test)
```

```
ConfusionMatrixDisplay.from_estimator(dec_tree, X_test, Y_test)
```



```
print("Accuracy on the test data ", metrics.accuracy_score(Y_test, prediction)*100, "%")
print("\nClassification report on the test data \n", metrics.classification_report(Y_test, prediction))
```

```
Accuracy on the test data 96.0 %
```


```
Classification report on the test data
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	17
1	0.95	0.95	0.95	19
2	0.93	0.93	0.93	14
accuracy			0.96	50
macro avg	0.96	0.96	0.96	50
weighted avg	0.96	0.96	0.96	50

Decision Tree (On Breast Cancer Dataset)

```
dec_tree1_c = DecisionTreeClassifier(criterion = 'entropy')
dec_tree1_c.fit(X_train_c, Y_train_c)
prediction_c = dec_tree1_c.predict(X_test_c)

ConfusionMatrixDisplay.from_estimator(dec_tree1, X_test_c, Y_test_c)
```



```
print("Accuracy on the test data ", metrics.accuracy_score(Y_test_c, prediction_c)*100, "%")
print("\nClassification report on the test data \n", metrics.classification_report(Y_test_c, prediction_c))
```

Accuracy on the test data 91.05263157894737 %

Classification report on the test data

	precision	recall	f1-score	support
--	-----------	--------	----------	---------


0	0.95	0.91	0.93	124
1	0.85	0.91	0.88	66

accuracy			0.91	190
macro avg	0.90	0.91	0.90	190
weighted avg	0.91	0.91	0.91	190

KNN (On Iris Dataset)

```
knn = KNeighborsClassifier()
knn.fit(X_train, Y_train)
prediction2 = knn.predict(X_test)

ConfusionMatrixDisplay.from_estimator(knn, X_test, Y_test)
```

```
print("Accuracy on the test data ", metrics.accuracy_score(Y_test, prediction2)*100, "%")
print("\nClassification report on the test data \n", metrics.classification_report(Y_test, prediction2))
```

Accuracy on the test data 96.0 %


Classification report on the test data

	precision	recall	f1-score	support
0	1.00	1.00	1.00	17
1	0.95	0.95	0.95	19
2	0.93	0.93	0.93	14
accuracy			0.96	50
macro avg	0.96	0.96	0.96	50
weighted avg	0.96	0.96	0.96	50

KNN (On Breast Cancer Dataset)

```
knn_c = KNeighborsClassifier()
knn_c.fit(X_train_c, Y_train_c)
prediction2_c = knn_c.predict(X_test_c)
```

```
ConfusionMatrixDisplay.from_estimator(knn_c, X_test_c, Y_test_c)
```



```
print("Accuracy on the test data ", metrics.accuracy_score(Y_test_c, prediction2_c)*100, "%")
print("\nClassification report on the test data \n", metrics.classification_report(Y_test_c, prediction2_c))
```

Accuracy on the test data 96.3157894736842 %


Classification report on the test data

	precision	recall	f1-score	support
0	0.96	0.98	0.97	124
1	0.97	0.92	0.95	66
accuracy			0.96	190
macro avg	0.96	0.95	0.96	190
weighted avg	0.96	0.96	0.96	190

Naïve Bayes (On Iris Dataset)

```
nb = GaussianNB()
nb.fit(X_train,Y_train)
prediction_nb=nb.predict(X_test)
```

```
ConfusionMatrixDisplay.from_estimator(nb, X_test, Y_test)
```



```
print("Accuracy on the test data ")
nb_score=metrics.accuracy_score(Y_test,prediction_nb)
print(f"Naive-bayes accuracy: {nb_score * 100}%")

print("\nClassification report on the test data ", metrics.classification_report(Y_test, prediction_nb))
```

Accuracy on the test data

Naive-bayes accuracy: 94.0%


Classification report on the test data

	precision	recall	f1-score	support
0	1.00	1.00	1.00	17
1	0.94	0.89	0.92	19
2	0.87	0.93	0.90	14
accuracy			0.94	50
macro avg	0.94	0.94	0.94	50
weighted avg	0.94	0.94	0.94	50

Naïve Bayes (On Breast Cancer Dataset)

```
nb_c = GaussianNB()
nb_c.fit(X_train_c,Y_train_c)
prediction_nbc=nb_c.predict(X_test_c)

ConfusionMatrixDisplay.from_estimator(nb_c, X_test_c, Y_test_c)
```



```
print("Accuracy on the test data ")
nb_score=metrics.accuracy_score(Y_test_c,prediction_nbc)
print(f"Naive-bayes accuracy: {nb_score * 100}%")

print("\nClassification report on the test data ", metrics.classification_report(Y_test_c, prediction_nbc))
```

Accuracy on the test data
Naive-bayes accuracy: 93.6842105263158%

					precision	recall	f1-score	support
	0	0.96	0.94	0.95	124			
	1	0.90	0.92	0.91	66			
	accuracy			0.94	190			
	macro avg	0.93	0.93	0.93	190			
	weighted avg	0.94	0.94	0.94	190			

Random Subsampling

Decision Tree (On Iris dataset)

```
acc = []

dec_tree = DecisionTreeClassifier(criterion = 'entropy')

for i in range(10):
    X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.33, random_state = i+1)

    dec_tree.fit(X_train, Y_train)
    prediction = dec_tree.predict(X_test)
    acc.append(metrics.accuracy_score(Y_test, prediction))

acc = sum(acc) / len(acc)

print("Accuracy on the test data ", acc*100, "%")

Accuracy on the test data 94.60000000000001 %
```

Decision Tree (On Breast Cancer dataset)

```
acc = []

dec_tree_c = DecisionTreeClassifier(criterion = 'entropy')

for i in range(10):
    X_train_c,X_test_c,Y_train_c,Y_test_c = train_test_split(X_c,Y_c,test_size=0.33, random_state = i+1)

    dec_tree_c.fit(X_train_c, Y_train_c)
    prediction_c = dec_tree_c.predict(X_test_c)
    acc.append(metrics.accuracy_score(Y_test_c, prediction_c))

acc = sum(acc) / len(acc)

print("Accuracy on the test data ", acc*100, "%")

Accuracy on the test data  94.04255319148938 %
```

KNN (On Iris Dataset)

```
acc = []

knn = KNeighborsClassifier()

for i in range(10):
    X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.33, random_state = i+1)
    knn.fit(X_train, Y_train)
    prediction2 = knn.predict(X_test)
    acc.append(metrics.accuracy_score(Y_test, prediction2))

acc = sum(acc) / len(acc)

print("Accuracy on the test data ", acc*100, "%")

Accuracy on the test data  95.19999999999997 %
```

KNN (On Breast Cancer Dataset)

```
acc = []

knn_c = KNeighborsClassifier()

for i in range(10):
    X_train_c,X_test_c,Y_train_c,Y_test_c = train_test_split(X_c,Y_c,test_size=0.33, random_state = i+1)
    knn_c.fit(X_train_c, Y_train_c)
    prediction2_c = knn_c.predict(X_test_c)
    acc.append(metrics.accuracy_score(Y_test_c, prediction2_c))

acc = sum(acc) / len(acc)

print("Accuracy on the test data ", acc*100, "%")

Accuracy on the test data  96.75531914893615 %
```

Naïve Bayes (On Iris Dataset)

```

acc = []

nb = GaussianNB()

for i in range(10):
    X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.33, random_state = i+1)
    nb.fit(X_train, Y_train)
    prediction_nb=nb.predict(X_test)
    acc.append(metrics.accuracy_score(Y_test, prediction_nb))

acc = sum(acc) / len(acc)

print("Accuracy on the test data ", acc*100, "%")

Accuracy on the test data  94.80000000000001 %

```

Naïve Bayes (On Breast Cancer Dataset)

```

acc = []

nb_c = GaussianNB()

for i in range(10):
    X_train_c,X_test_c,Y_train_c,Y_test_c = train_test_split(X_c,Y_c,test_size=0.33, random_state = i+1)
    nb_c.fit(X_train_c, Y_train_c)
    prediction_nbc=nb_c.predict(X_test_c)
    acc.append(metrics.accuracy_score(Y_test_c, prediction_nbc))

acc = sum(acc) / len(acc)

print("Accuracy on the test data ", acc*100, "%")

Accuracy on the test data  93.88297872340428 %

```

Cross Validation

Decision Tree (On Iris Dataset)

```

[121] from sklearn.model_selection import KFold
      from sklearn.model_selection import cross_val_score

IRIS DATASET

[122] dec_tree = DecisionTreeClassifier(criterion = 'entropy')

      score1 = cross_val_score(dec_tree, X, Y, cv=4)

[123] acc_cv_iris_dt = score1.mean()*100

      print("Accuracy on the test data ", acc_cv_iris_dt, "%")

Accuracy on the test data  95.99928876244665 %

```

Decision Tree (On Breast Cancer Dataset)

```

dec_tree_c = DecisionTreeClassifier(criterion = 'entropy')

score1 = cross_val_score(dec_tree_c, X_c, Y_c, cv=4)

acc_cv_breast_dt = score1.mean()*100

print("Accuracy on the test data ", acc_cv_breast_dt, "%")

Accuracy on the test data  91.91248891953117 %

```

KNN (On Iris and Breast Cancer Dataset)

```

IRIS DATASET

[126] knn = KNeighborsClassifier()

      score2 = cross_val_score(knn, X, Y, cv=4)

[127] acc_cv_iris_knn = score2.mean()*100

      print("Accuracy on the test data ", acc_cv_iris_knn, "%")

      Accuracy on the test data  94.66571834992888 %

BREAST CANCER DATASET

[128] knn_C = KNeighborsClassifier()

      score2 = cross_val_score(knn_C, X_c, Y_c, cv=4)

[129] acc_cv_breast_knn = score2.mean()*100

      print("Accuracy on the test data ", score2.mean()*100, "%")

      Accuracy on the test data  96.48626021865458 %

```

Naïve bayes (On Iris and Breast Cancer Dataset)

```

IRIS DATASET

[130] nb = GaussianNB()

      score3 = cross_val_score(nb, X, Y, cv=4)

[131] acc_cv_iris_nb = score3.mean()*100

      print("Accuracy on the test data ", acc_cv_iris_nb, "%")

      Accuracy on the test data  95.34139402560456 %

BREAST CANCER DATASET

[132] nb_C = GaussianNB()

      score3 = cross_val_score(nb, X_c, Y_c, cv=4)

[133] acc_cv_breast_nb = score3.mean()*100

      print("Accuracy on the test data ", score3.mean()*100, "%")

      Accuracy on the test data  92.97498276371515 %

```

Result:

Holdout

```
▶ print("Accuracy of decision tree model on iris data : ", acc_hold_iris1)
print("Accuracy of decision tree model on breast data : ", acc_hold_breast1)

print("Accuracy of KNN model on iris data : ", acc_hold_iris_knn )
print("Accuracy of KNN model on iris data : ", acc_hold_breast_knn )

print("Accuracy of Naive Bayes model on iris data : ", acc_hold_iris_nb )
print("Accuracy of Naive Bayes model on iris data : ", acc_hold_breast_nb)
```

```
☐ Accuracy of decision tree model on iris data : 97.36842105263158
Accuracy of decision tree model on breast data : 95.1048951048951
Accuracy of KNN model on iris data : 97.36842105263158
Accuracy of KNN model on iris data : 95.1048951048951
Accuracy of Naive Bayes model on iris data : 97.36842105263158
Accuracy of Naive Bayes model on iris data : 94.4055944055944
```

Random Subsampling

```
[135] print("Accuracy of decision tree model on iris data : ", acc_rs_iris_dt)
print("Accuracy of decision tree model on breast data : ", acc_rs_breast_dt)

print("Accuracy of KNN model on iris data : ", acc_rs_iris_knn )
print("Accuracy of KNN model on iris data : ", acc_rs_breast_knn )

print("Accuracy of Naive Bayes model on iris data : ", acc_rs_iris_nb )
print("Accuracy of Naive Bayes model on iris data : ", acc_rs_breast_nb)
```

```
Accuracy of decision tree model on iris data : 94.73684210526315
Accuracy of decision tree model on breast data : 92.72727272727272
Accuracy of KNN model on iris data : 95.26315789473682
Accuracy of KNN model on iris data : 96.50349650349652
Accuracy of Naive Bayes model on iris data : 94.73684210526315
Accuracy of Naive Bayes model on iris data : 93.77622377622379
```

Cross validation

```
▶ print("Accuracy of decision tree model on iris data : ", acc_cv_iris_dt)
print("Accuracy of decision tree model on breast data : ", acc_cv_breast_dt)

print("Accuracy of KNN model on iris data : ", acc_cv_iris_knn )
print("Accuracy of KNN model on iris data : ", acc_cv_breast_knn )

print("Accuracy of Naive Bayes model on iris data : ", acc_cv_iris_nb )
print("Accuracy of Naive Bayes model on iris data : ", acc_cv_breast_nb)
```

```
Accuracy of decision tree model on iris data : 95.99928876244665
Accuracy of decision tree model on breast data : 91.91248891953117
Accuracy of KNN model on iris data : 94.66571834992888
Accuracy of KNN model on iris data : 96.48626021865458
Accuracy of Naive Bayes model on iris data : 95.34139402560456
Accuracy of Naive Bayes model on iris data : 92.97498276371515
```


```
import matplotlib.pyplot as plt

X = ['Decision Tree', 'KNN', 'Naive Bayes']
Yiris = [acc_hold_iris1, acc_hold_iris_knn, acc_hold_iris_nb]
Ybc = [acc_hold_breast1, acc_hold_breast_knn, acc_hold_breast_nb]

X_axis = np.arange(len(X))

plt.bar(X_axis - 0.2, Yiris, 0.4, label="Iris")
plt.bar(X_axis + 0.2, Ybc, 0.4, label="Breast Cancer")

plt.xlabel("Classifiers Used")
plt.ylabel("% of accuracy achieved")
plt.title("Accuracy of Classifiers using holdout")
plt.xticks(X_axis, X, rotation = 20)
plt.legend()
plt.show()
```




```
X = ['Decision Tree', 'KNN', 'Naive Bayes']
Yiris = [acc_rs_iris_dt, acc_rs_iris_knn, acc_rs_iris_nb]
Ybc = [acc_rs_breast_dt, acc_rs_breast_knn, acc_rs_breast_nb]

X_axis = np.arange(len(X))

plt.bar(X_axis - 0.2, Yiris, 0.4, label="Iris")
plt.bar(X_axis + 0.2, Ybc, 0.4, label="Breast Cancer")

plt.xlabel("Classifiers Used")
plt.ylabel("% of accuracy achieved")
plt.title("Accuracy of Classifiers using Random Subsampling")
plt.xticks(X_axis, X, rotation = 20)
plt.legend()
plt.show()
```





```
X = ['Decision Tree', 'KNN', 'Naive Bayes']
Yiris = [acc_cv_iris_dt, acc_cv_iris_knn, acc_cv_iris_nb]
Ybc = [acc_cv_breast_dt, acc_cv_breast_knn, acc_cv_breast_nb]

X_axis = np.arange(len(X))

plt.bar(X_axis - 0.2, Yiris, 0.4, label="Iris")
plt.bar(X_axis + 0.2, Ybc, 0.4, label="Breast Cancer")

plt.xlabel("Classifiers Used")
plt.ylabel("% of accuracy achieved")
plt.title("Accuracy of Classifiers using Cross Validation")
plt.xticks(X_axis, X, rotation = 20)
plt.legend()
plt.show()
```



Q6. Use Simple Kmeans, DBScan, Hierarchical clustering algorithms for clustering. Compare the performance of clusters by changing the parameters involved in the algorithms.

```
import pandas as pd
import numpy as np
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn import metrics
```

```
iris = load_iris()
```

```
df_iris = pd.DataFrame(iris.data, columns=iris.feature_names)
```

```
df_iris['Class'] = iris.target
```

```
df_iris.head()
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	Class
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0

```
df_iris.shape
```

```
(150, 5)
```

```
df_iris.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   sepal length (cm)      150 non-null   float64
1   sepal width (cm)       150 non-null   float64
2   petal length (cm)      150 non-null   float64
3   petal width (cm)       150 non-null   float64
4   Class                  150 non-null   int64
dtypes: float64(4), int64(1)
memory usage: 6.0 KB
```

```
df_iris.describe()
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	Class
count	150.000000	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333	1.000000
std	0.828066	0.435866	1.765298	0.762238	0.819232
min	4.300000	2.000000	1.000000	0.100000	0.000000
25%	5.100000	2.800000	1.600000	0.300000	0.000000
50%	5.800000	3.000000	4.350000	1.300000	1.000000
75%	6.400000	3.300000	5.100000	1.800000	2.000000
max	7.900000	4.400000	6.900000	2.500000	2.000000

```

scaler = StandardScaler()

scaled_X = scaler.fit_transform(df_iris.iloc[:, :-1])

df_iris_scaled = pd.DataFrame(scaled_X, columns = list(df_iris.columns[:-1]))

df_iris_scaled["class"] = df_iris["Class"]

df_iris_scaled.head()

```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	class
0	-0.900681	1.019004	-1.340227	-1.315444	0
1	-1.143017	-0.131979	-1.340227	-1.315444	0
2	-1.385353	0.328414	-1.397064	-1.315444	0
3	-1.506521	0.098217	-1.283389	-1.315444	0
4	-1.021849	1.249201	-1.340227	-1.315444	0

```
df_iris_scaled.describe()
```


	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	class
count	1.500000e+02	1.500000e+02	1.500000e+02	1.500000e+02	150.000000
mean	-1.468455e-15	-1.823726e-15	-1.610564e-15	-9.473903e-16	1.000000
std	1.003350e+00	1.003350e+00	1.003350e+00	1.003350e+00	0.819232
min	-1.870024e+00	-2.433947e+00	-1.567576e+00	-1.447076e+00	0.000000
25%	-9.006812e-01	-5.923730e-01	-1.226552e+00	-1.183812e+00	0.000000
50%	-5.250608e-02	-1.319795e-01	3.364776e-01	1.325097e-01	1.000000
75%	6.745011e-01	5.586108e-01	7.627583e-01	7.906707e-01	2.000000
max	2.492019e+00	3.090775e+00	1.785832e+00	1.712096e+00	2.000000

```

import matplotlib.pyplot as plt

plt.scatter(df_iris_scaled["petal length (cm)"], df_iris_scaled["petal width (cm)"])

```



```
X = df_iris_scaled.iloc[:, :-1]
Y = df_iris_scaled["class"]
```

Kmeans Clustering

```
from sklearn.cluster import KMeans
```

```
kmeans = KMeans(n_clusters = 3, random_state=1, n_init = 'auto')
kmeans.fit(X)
```

```
▼ KMeans
KMeans(n_clusters=3, n_init='auto', random_state=1)
```


```
center = kmeans.cluster_centers_
```

```
identified_clusters = kmeans.fit_predict(X)
identified_clusters
```

```
temp1 = X[identified_clusters == 0]
temp2 = X[identified_clusters == 1]
temp3 = X[identified_clusters == 2]
```

```
#Plotting the results
```

```
plt.scatter(temp1.iloc[:,2] , temp1.iloc[:,3] , color = 'red')
plt.scatter(temp2.iloc[:,2] , temp2.iloc[:,3] , color = 'black')
plt.scatter(temp3.iloc[:,2], temp3.iloc[:, 3], color='blue')
plt.show()
```




DBScan

```
from sklearn.cluster import DBSCAN
from sklearn.neighbors import NearestNeighbors
```

```
neigh = NearestNeighbors(n_neighbors=2)
nbrs = neigh.fit(X)
distances, indices = nbrs.kneighbors(X)
```

```
distances = np.sort(distances, axis=0)
distances = distances[:,1]
plt.plot(distances)
plt.grid()
plt.show()
```



```
dbscan = DBSCAN(eps=0.6, min_samples=3)
dbscan.fit(X)
```

```
DBSCAN
DBSCAN(eps=0.6, min_samples=3)
```

```
identified_clusters_db = dbscan.fit_predict(X)
identified_clusters_db
```

```
unique_clusters = np.unique(identified_clusters_db)
```

```
li = []
```


```
for i in range(len(unique_clusters)):
    li.append(X[identified_clusters_db == unique_clusters[i]])
```

```
import random
no_of_colors = len(li)
```

```
color=["#"+''.join([random.choice('0123456789ABCDEF') for i in range(6)]) for j in range(no_of_colors)]
```

```
#Plotting the results
```

```
for i in range(len(li)):
    plt.scatter(li[i].iloc[:,2] , li[i].iloc[:,3] , color = color[i])
plt.show()
```



Hierarchical Clustering

```
from sklearn.cluster import AgglomerativeClustering
```


```
hier_clus = AgglomerativeClustering(n_clusters=3)
hier_clus.fit(X)
```

```
identified_clusters_hc = hier_clus.fit_predict(X)
identified_clusters_hc
```

```
temp1 = X[identified_clusters_hc == 0]
temp2 = X[identified_clusters_hc == 1]
temp3 = X[identified_clusters_hc == 2]
```

```
#Plotting the results
```

```
plt.scatter(temp1.iloc[:,2] , temp1.iloc[:,3] , color = 'red')
plt.scatter(temp2.iloc[:,2] , temp2.iloc[:,3] , color = 'black')
plt.scatter(temp3.iloc[:,2], temp3.iloc[:, 3], color='blue')
plt.show()
```



```
from scipy.cluster import hierarchy

clusters = hierarchy.linkage(X, method="ward")

plt.figure(figsize=(8, 6))
dendrogram = hierarchy.dendrogram(clusters)
```

